

# Predicting from Strings:

## Language Model Embeddings for Bayesian Optimization

Yilong Chen

Department of Statistics, University of Chicago

Nov, 2024



- 1 Preliminaries
- 2 Regressor
- 3 Pretraining and Inference
- 4 Attention Mechanism
- 5 Experiment

# 1 Preliminaries

## 2 Regressor

## 3 Pretraining and Inference

## 4 Attention Mechanism

## 5 Experiment

# Blackbox Optimization Problem

**Goal:** Maximize a real-valued function  $f$  over a search space  $\mathcal{X}$ :

$$x^* = \arg \max_{x \in \mathcal{X}} f(x) \quad (1)$$

## Challenges:

- $f$  is a blackbox function, i.e., no analytical expression or gradients.
- Evaluation of  $f$  is expensive or limited.
- Requires balance between **exploration** and **exploitation**.

# Definition of a Regressor

**Regressor:** A prediction model that outputs the distribution of  $f(\cdot)$  values over a query point  $x$ , given the history of evaluations:

$$\{(x_s, y_s)\}_{s=1}^t$$

where  $y_s = f(x_s)$  for  $s = 1, 2, \dots, t$ .

## Key Properties:

- *Learnability:* The regressor can be trained using offline or online data.
- *Distributional Output:* Provides mean and uncertainty estimates for predictions.

# Acquisition Function for Exploration and Exploitation

**Definition:** The regressor is transformed into an acquisition function:

$$a_{t+1}(x) : \mathcal{X} \rightarrow \mathbb{R}$$

which guides the optimization process.

**Next Proposal:**

$$x_{t+1} = \arg \max_{x \in \mathcal{X}} a_{t+1}(x) \quad (2)$$

**Acquisition Optimizer:**

- Samples  $x \in \mathcal{X}$  efficiently, using zeroth-order or evolutionary algorithms.
- Balances exploration (high uncertainty) and exploitation (high mean).

# Bayesian Optimization Loop

## Procedure:

- 1 Initialize a set of observations  $\{(x_s, y_s)\}_{s=1}^t$ .
- 2 Train a regressor on the given data.
- 3 Compute the acquisition function  $a_{t+1}(x)$  based on the regressor's predictions.
- 4 Find the next proposal:

$$x_{t+1} = \arg \max_{x \in \mathcal{X}} a_{t+1}(x)$$

- 5 Evaluate  $f(x_{t+1})$  and update the history.
- 6 Repeat until a stopping criterion is met.

**Outcome:** Approximate the global maximum of  $f(x)$ .

- 1 Preliminaries
- 2 Regressor
- 3 Pretraining and Inference
- 4 Attention Mechanism
- 5 Experiment



# Embedding-Based Regressor: Concept

**Purpose:** Convert input  $x \in \mathcal{X}$  into a fixed-length representation for regression.

## Embedding Process:

- The embedder  $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$  maps  $x$  to a vector  $\bar{x} \in \mathbb{R}^d$ .
- Input  $x$  is represented as a **string**, which is processed by a language model (e.g., T5 encoder).
- A forward pass through the model produces token logits in  $\mathbb{R}^{L \times d}$ .

**Insert Graph:** Diagram of embedding process:

- Input string  $\rightarrow$  Tokenization  $\rightarrow$  Language Model  $\rightarrow$  Average Pooling  $\rightarrow \bar{x}$ .

# Embedding-Based Regressor: Final Representation

## Final Step:

- Perform **average pooling** across the token axis to generate a fixed-length embedding in  $\mathbb{R}^d$ .
- Output:  $\bar{x} \in \mathbb{R}^d$ .

## Benefits:

- Handles arbitrary input types (e.g., JSON, categorical).
- Generates consistent, fixed-length feature vectors for regression tasks.

# In-context Regression Transformer: Input Sequence

## Input Sequence:

$$(\bar{x}_1 \oplus \bar{y}_1), \dots, (\bar{x}_t \oplus \bar{y}_t) \quad (3)$$

where:


- $\bar{x}_i \in \mathbb{R}^d$ : Embedding of trial  $x_i$ .
- $\bar{y}_i \in \mathbb{R}^d$ : Feature representation of observed value  $y_i$ , obtained via a trainable projection.
- $\oplus$ : Concatenation operator.

# In-context Regression Transformer: Prediction

## Prediction Process:

- Append query  $(\bar{x} \oplus \bar{0})$ , where  $\bar{0}$  is a placeholder.
- Perform a forward pass through the Transformer.
- The output feature corresponding to  $t + 1$  predicts:

$$\mathcal{N}(\mu_{t+1}(x), \sigma_{t+1}^2(x)) \quad (4)$$



# Additional Techniques

## Parallel Predictions:

- Simultaneously predict over a set of  $k$  target points:

$$(\bar{x}_{t+1} \oplus \bar{0}), \dots, (\bar{x}_{t+k} \oplus \bar{0}) \quad (5)$$

- Custom attention pattern allows tokens to attend to history but not to targets.

# Additional Techniques

## y-Normalization:

- Shift objectives to have zero mean and scale by standard deviation.
- Transform  $y$  values into  $[0, 1]$  for numerical stability.

## Encoding Metadata:

- Include task-specific metadata  $m$  as part of input:

$$\bar{x} \rightarrow (\bar{x} \oplus m) \quad (6)$$

- Useful for providing additional context about the objective or search space.

- 1 Preliminaries
- 2 Regressor
- 3 Pretraining and Inference**
- 4 Attention Mechanism
- 5 Experiment

# Pretraining and Inference: Overview

## Task Definition:

- A task  $\mathcal{T} = (f, \mathcal{X})$  represents a specific objective function  $f$  over a search space  $\mathcal{X}$ .

## Pretraining:

- Assumes a collection of offline **training tasks**  $\{\mathcal{T}_1, \mathcal{T}_2, \dots\}$ .
- Each task has its own evaluated trials  $\{(x_s, y_s)\}_{s=1}^T$ .
- $T$ : Offline trajectory length (task-specific).

## Purpose:

- Learn from historical offline data to generalize across tasks.



# Pretraining: Loss Function

## Frozen Embedder:

- During pretraining, the weights  $\theta$  of the ICL regression Transformer are optimized.
- The embedder remains **frozen**.

## Training Example:

- History:  $\{(x_s, y_s)\}_{s=1}^{t'}$  for some  $t' \in [0, T]$ .
- Targets:  $\{(x_{t'+i}, y_{t'+i})\}_{i=1}^{T-t'}$ .

## Loss Function:

$$\sum_{i=1}^{T-t'} \ell_{\theta}(x_{t'+i}, y_{t'+i}; \{(x_s, y_s)\}_{s=1}^{t'}), \quad (7)$$

where:

- $\ell_{\theta}$  is the negative log-likelihood of the Gaussian distribution.

# Inference: Overview

## At Inference:

- The Transformer predicts the mean  $\mu_{t+1}(x)$  and standard deviation  $\sigma_{t+1}(x)$  for the target.
- The acquisition function is defined as:

$$a_{t+1}(x) = \mu_{t+1}(x) + \sqrt{\beta} \cdot \sigma_{t+1}(x), \quad (8)$$

where  $\sqrt{\beta}$  is a problem-dependent constant.

## Optimization:

- Uses a zeroth-order optimizer (e.g., evolutionary search) to maximize  $a_{t+1}(x)$ .
- Requires only **forward passes** (no gradient-based optimization).

# Handling Distributional Shifts

## Challenges:

- Distributional shifts may occur between pretraining and inference.

## Solutions:

- **Data Augmentation:** Randomize parameter names during pretraining.
- **Search Space Transformation:** Adjust search space at inference to align with pretraining conditions.

- 1 Preliminaries
- 2 Regressor
- 3 Pretraining and Inference
- 4 Attention Mechanism**
- 5 Experiment

# Attention in Transformer for $\bar{x} \oplus \bar{0}$

**Goal:** Predict a distribution  $\mathcal{N}(\mu_{t+1}(\bar{x}), \sigma_{t+1}^2(\bar{x}))$  for a query point  $\bar{x} \oplus \bar{0}$ .

**Input Sequence:**

$$\mathcal{S} = \{(\bar{x}_1 \oplus \bar{y}_1), (\bar{x}_2 \oplus \bar{y}_2), \dots, (\bar{x}_t \oplus \bar{y}_t), (\bar{x} \oplus \bar{0})\}$$

**Attention Formula:**

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right) V$$

**Key Idea:** Query  $\bar{x} \oplus \bar{0}$  computes attention weights over historical points  $\{(\bar{x}_i \oplus \bar{y}_i)\}_{i=1}^t$ .

# How Attention Works for $\bar{x} \oplus \bar{0}$

## Steps:

- Compute attention weights:

$$\alpha_{t+1,i} = \text{softmax} \left( \frac{K(\bar{x} \oplus \bar{0}) \cdot Q(\bar{x}_i \oplus \bar{y}_i)}{\sqrt{d_k}} \right)$$

- Aggregate historical information:

$$z_{t+1} = \sum_{i=1}^t \alpha_{t+1,i} V_i(\bar{x}_i \oplus \bar{y}_i)$$

where  $V_i$  is the value vector of  $(\bar{x}_i \oplus \bar{y}_i)$ .

- Use  $z_{t+1}$  to predict  $\mu_{t+1}(\bar{x})$  and  $\sigma_{t+1}^2(\bar{x})$ .

## Output:

$$\mathcal{N}(\mu_{t+1}(\bar{x}), \sigma_{t+1}^2(\bar{x}))$$

- 1 Preliminaries
- 2 Regressor
- 3 Pretraining and Inference
- 4 Attention Mechanism
- 5 Experiment**

# Synthetic Optimization: Overview

## Definition:

- In common optimization scenarios, the search space is modeled as a flat Cartesian product of:
  - **Float parameters:** Continuous-valued parameters.
  - **Categorical parameters:** Discrete-valued parameters.
- The goal is to evaluate the performance of **Embed-then-Regress** on such tasks.

## Key Representation:

- Each  $x$  (input) is represented as a *JSON string*.
- Example:

`{"p0" : 0.3, "p1" : 4}`

- Where:
  - $p0$ : Continuous parameter.
  - $p1$ : Integer parameter.



# Benchmark: Blackbox Optimization Benchmarking (BBOB)

## Setup:

- Uses the **BBOB suite** (ElHara et al., 2019), one of the most widely used synthetic function benchmarks.
- Contains 24 objectives over **continuous search spaces**.

## Training-Test Split:

- Original functions are divided into **training** and **test sets**.
- Apply transformations (e.g., shifting, rotating, discretizing, etc.) to generate diverse functions.
- Induce non-continuous search spaces with categorical parameters.

**Goal:** Evaluate the generalization of the trained model on unseen tasks.

# Comparison with GP-Bandit

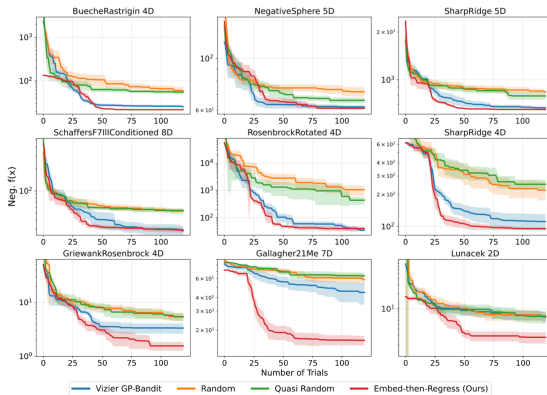
## Baseline:

- The traditional baseline is **GP-Bandit** (Song et al., 2024c), a UCB-based Bayesian Optimization method.
- Uses the same acquisition optimizer (**Firefly**) as Embed-then-Regress.

## Findings:

- Embed-then-Regress is generally **comparable** to GP-Bandit in performance.
- In some cases, it **significantly outperforms** GP-Bandit, especially on:
  - Non-continuous search spaces.
  - Test functions requiring task generalization.

# Performance



*Thank You*